

Quality-driven development of multi-cloud applications

Giuliano Casale

Department of Computing

Imperial College London

g.casale@imperial.ac.uk

Joint work with:

D. Dubois, R. Osman, J. F. Pérez

Computing@Imperial

Dept. of Computing:

- Languages & Systems
- Logic & AI
- Quantitative Analysis ←
- Software Engineering
- Vision & Robotics

My research:

- Cloud computing
- Performance engineering
- Stochastic modelling





- Large-Scale Integrated Project 9M€ (2012-2015)
- Quality-driven development for multi-cloud SW
- Tools: <http://www.modaclouds.eu/software>



Multi-Cloud Application

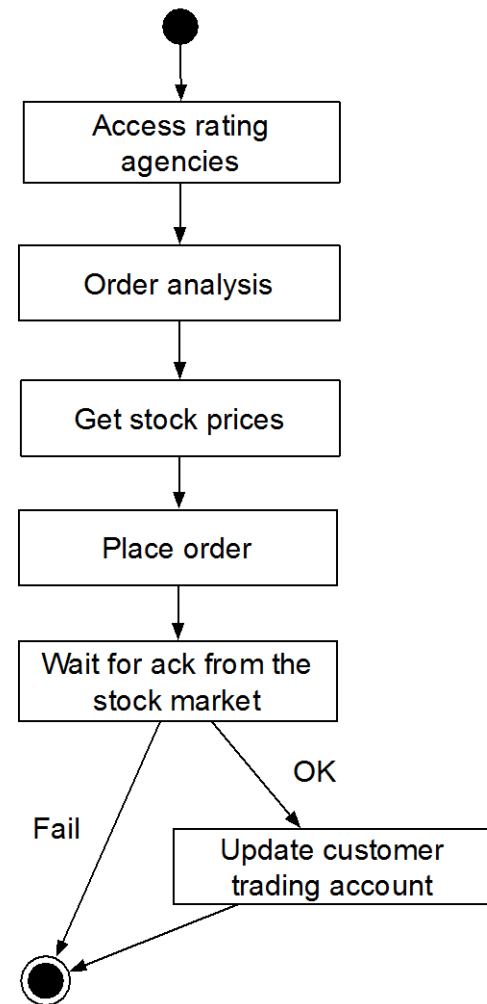
- The application can be deployed on different clouds



MDE for Multi-Cloud Software

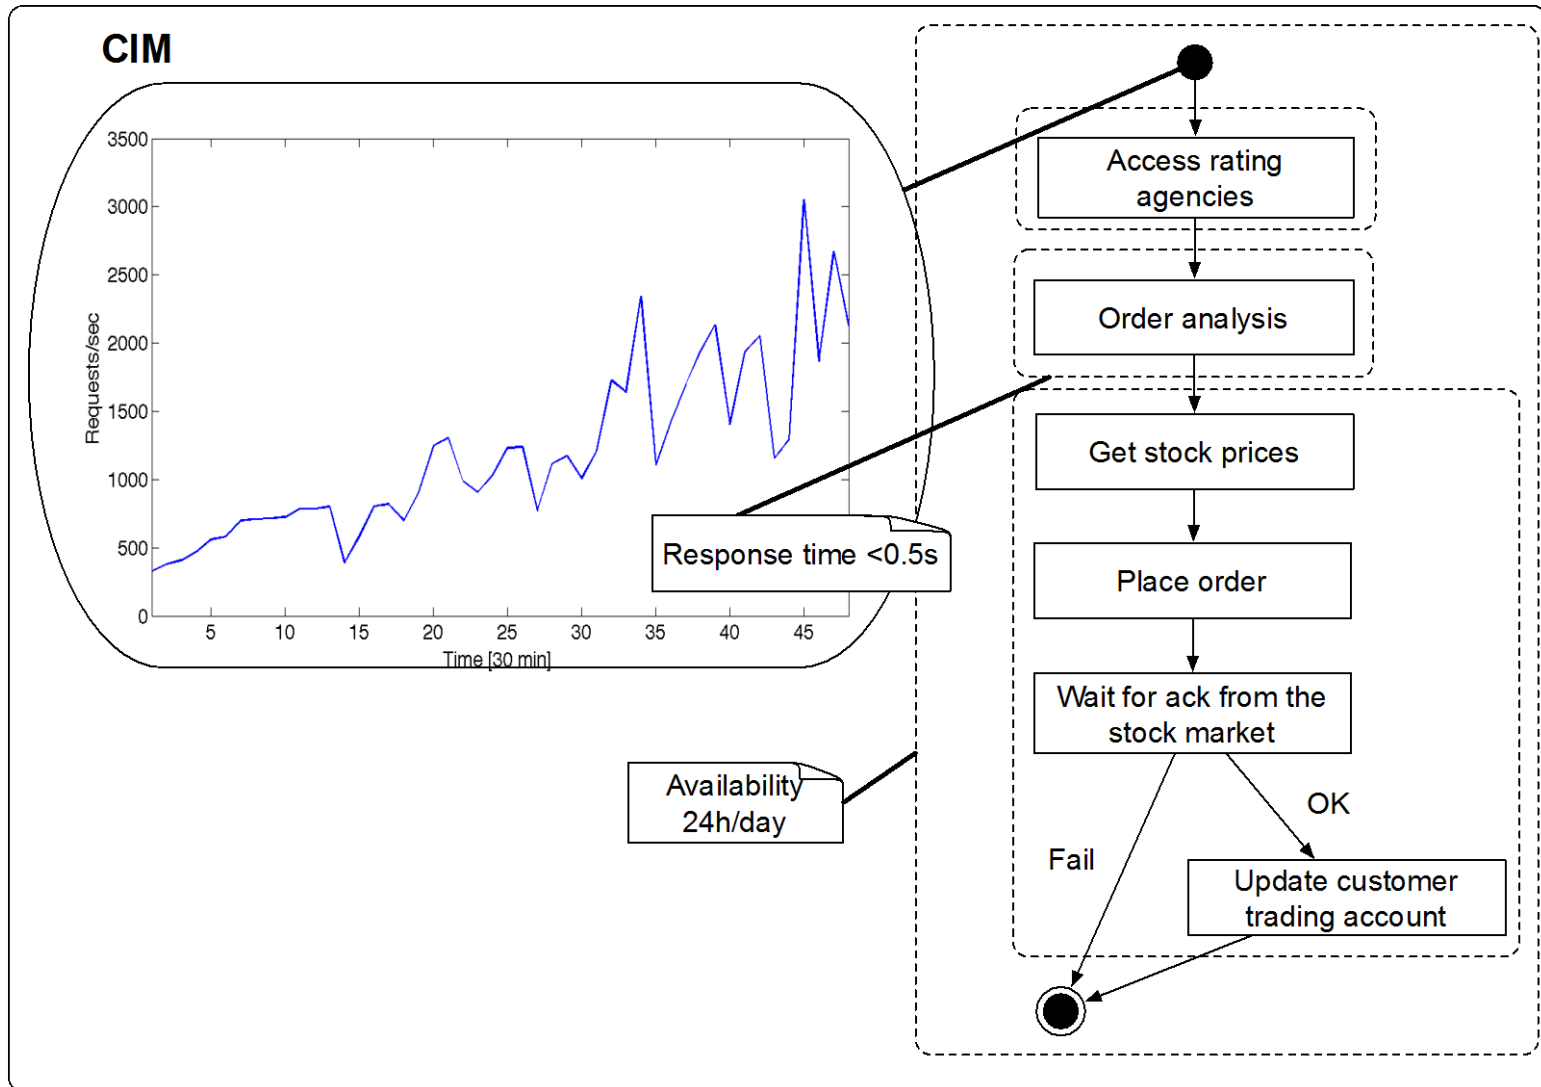
Functional Modelling

CCIM



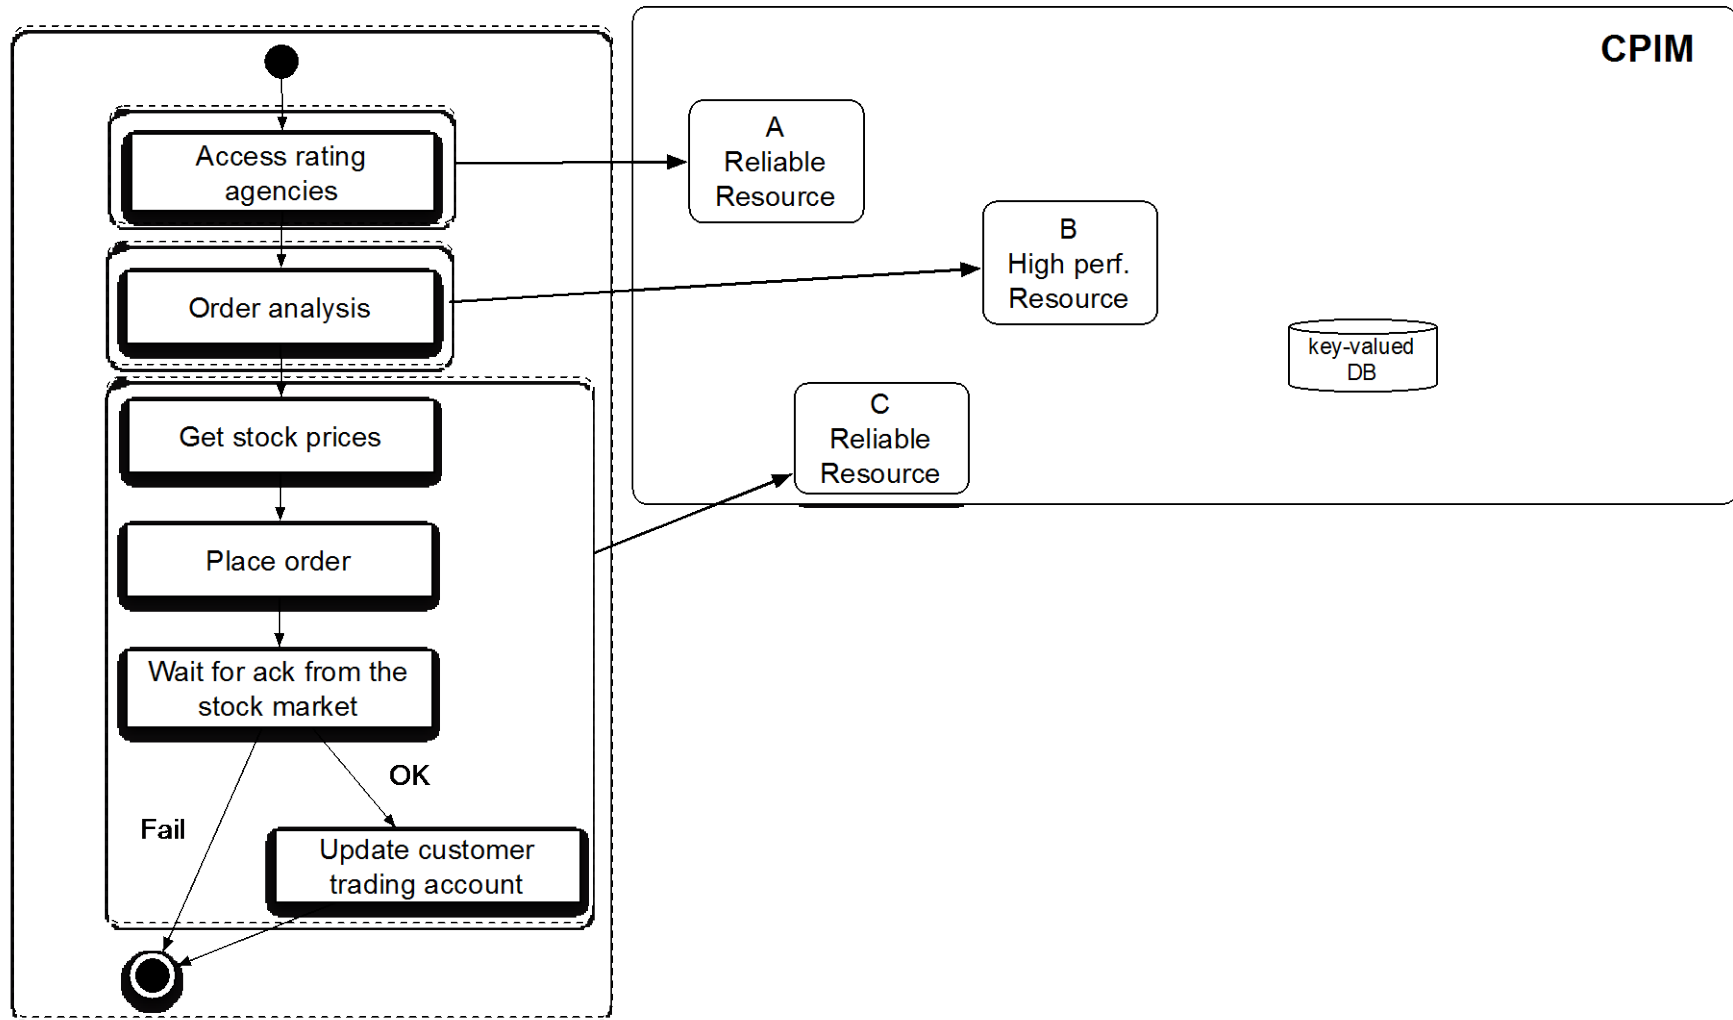
MDE for Multi-Cloud Software

QoS Requirements



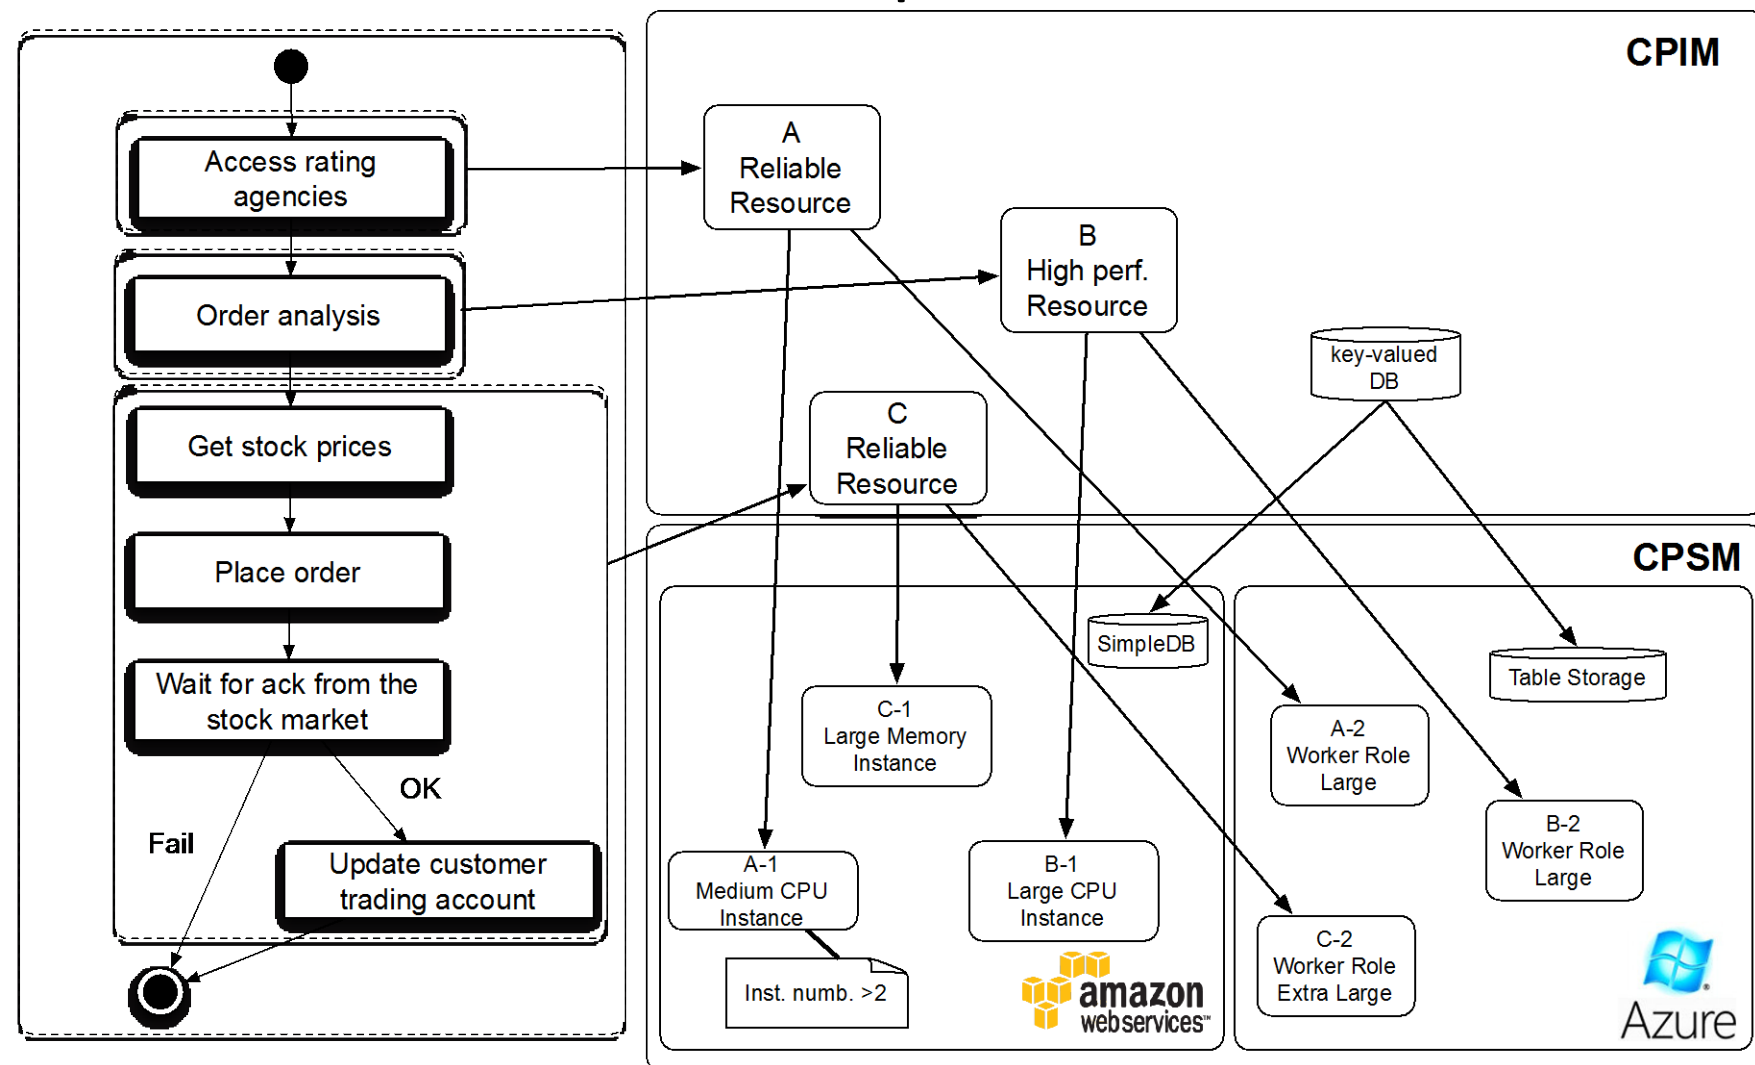
MDE for Multi-Cloud Software

Platform-Independent Models



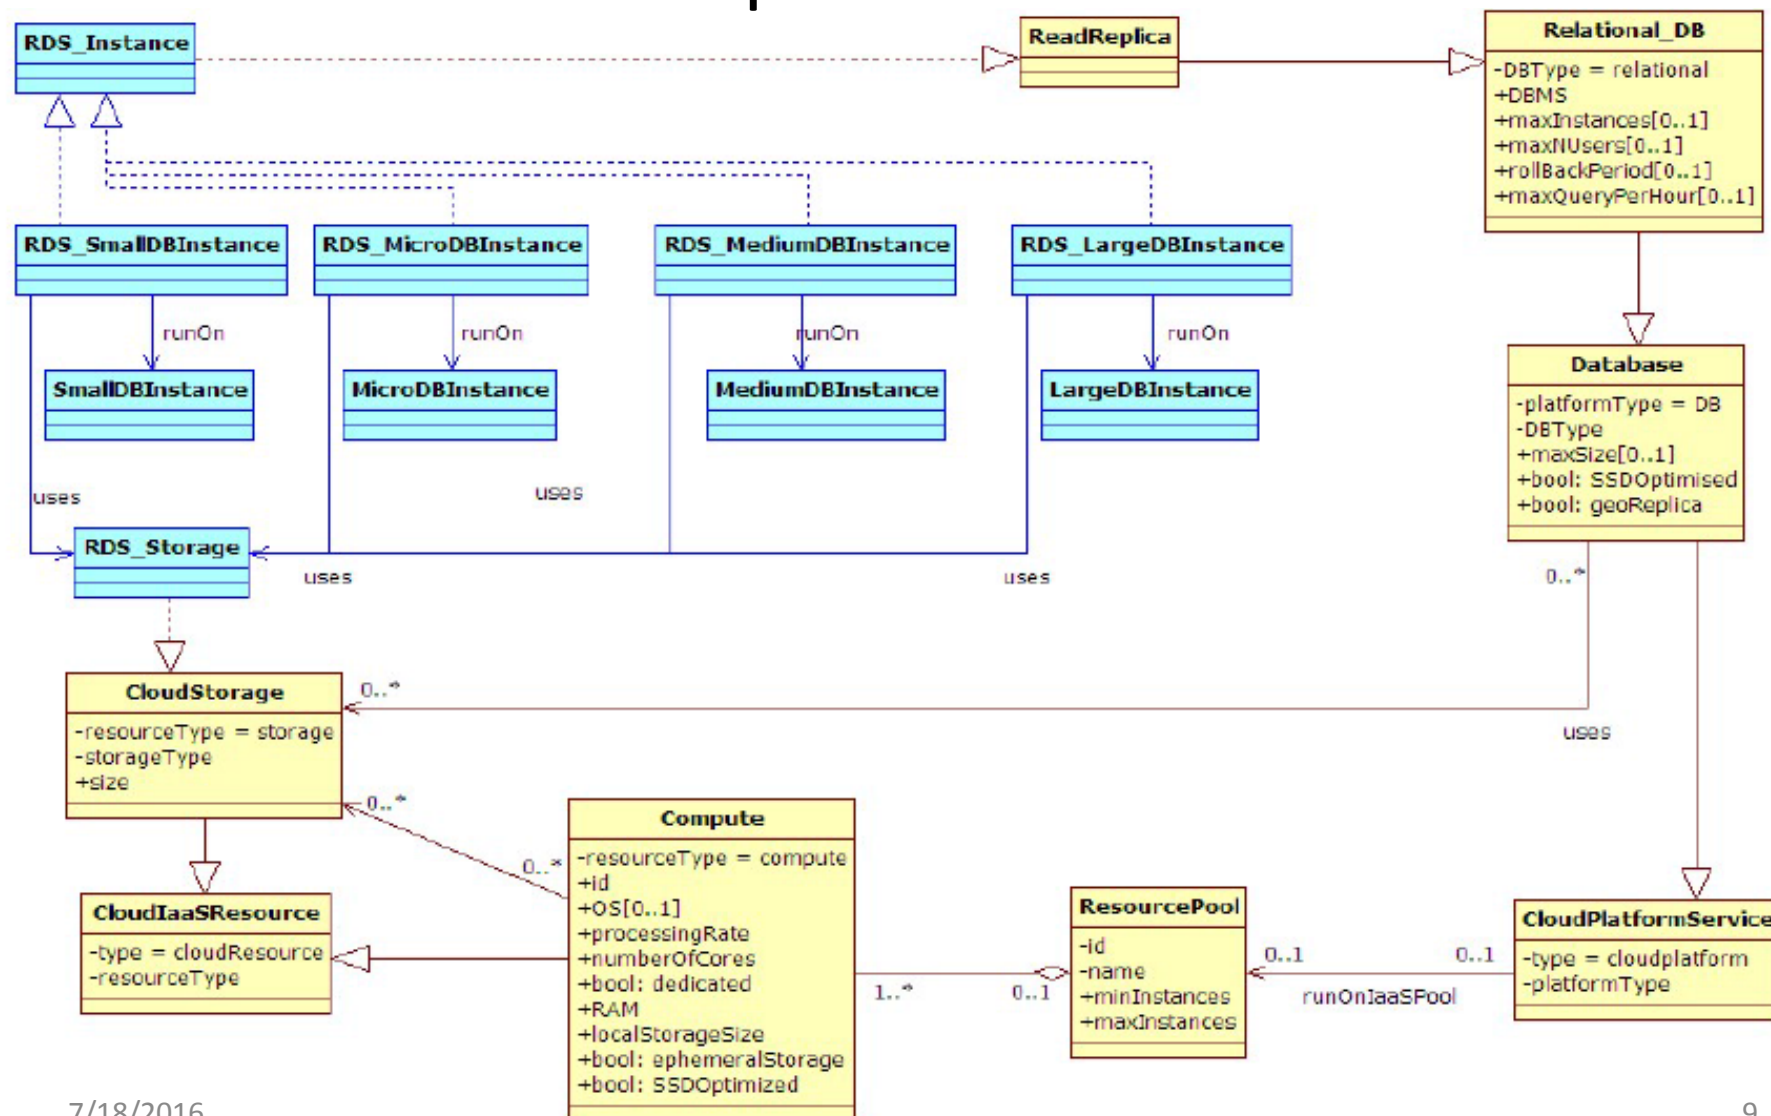
MDE for Multi-Cloud Software

Platform-Specific Models

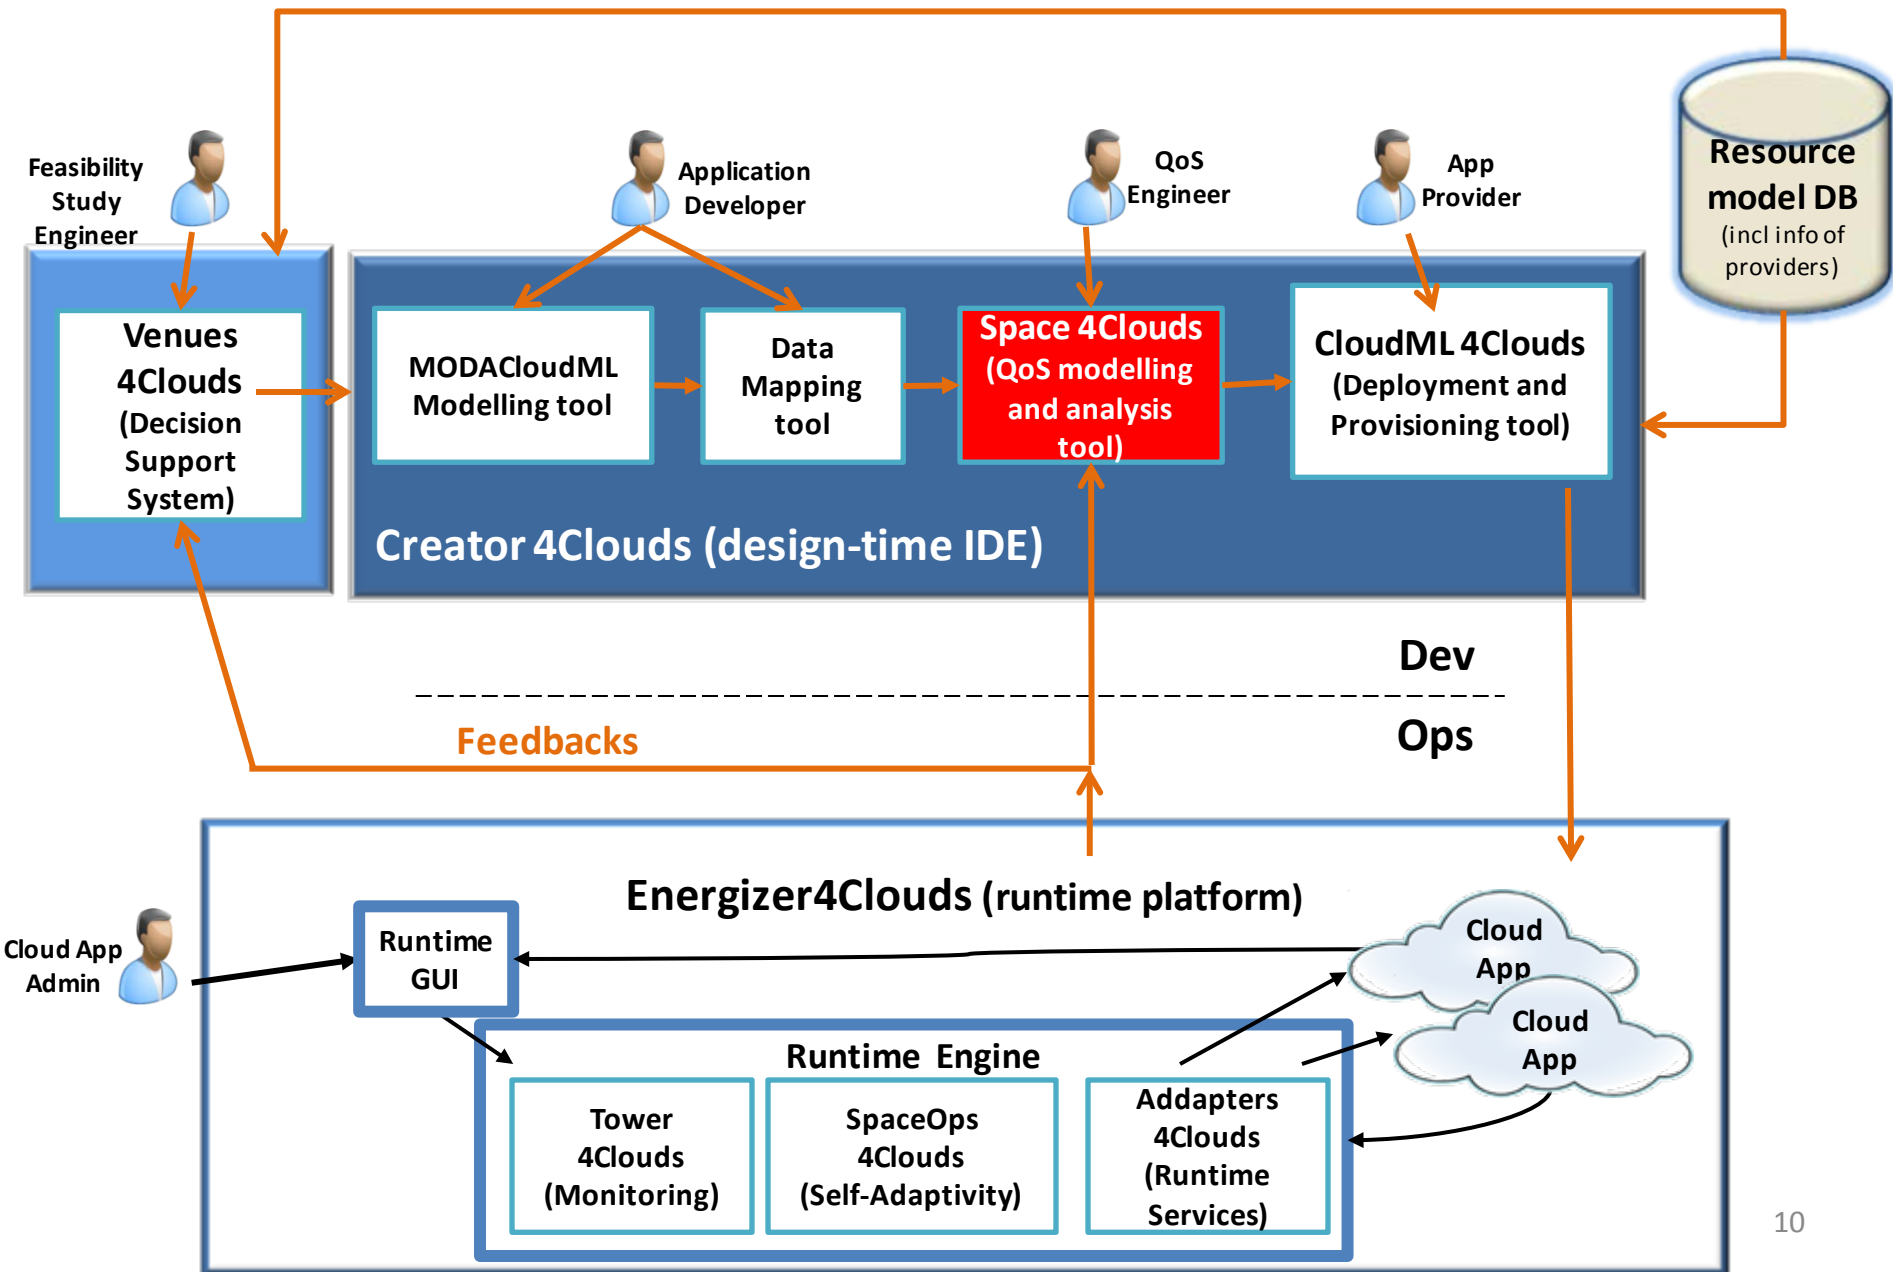


CloudML DSL (cloudml.org)

Example: Amazon RDS




MODAClouds Platform

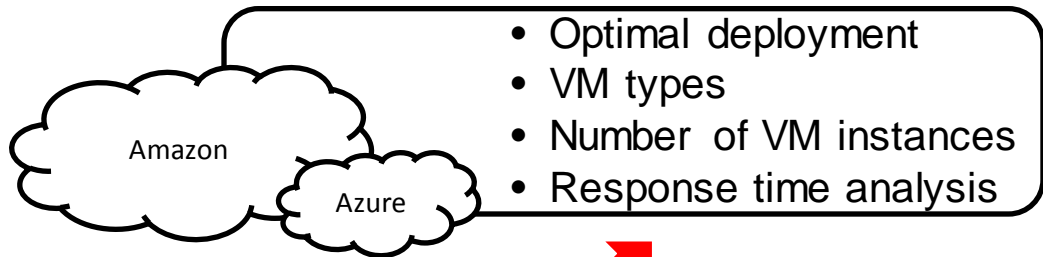


QoS Modelling and Analysis

Service and QoS model



Transformation
CloudML → QN




Output

Performance Analysis




- Performance prediction
- Optimization QoS
- Capacity planning



**Operational
Data**

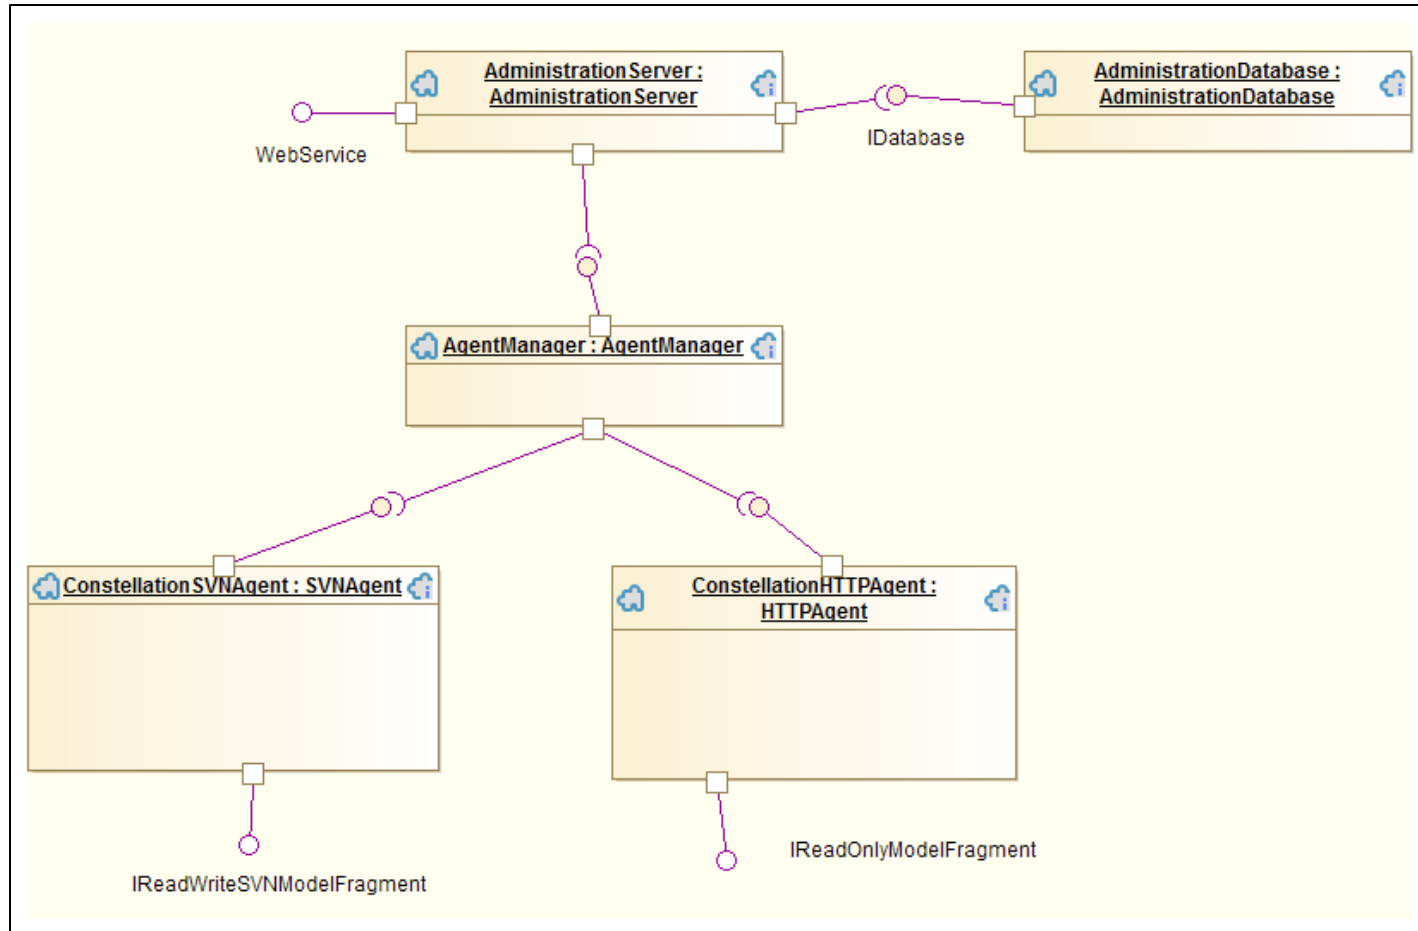


Performance Modelling Solvers



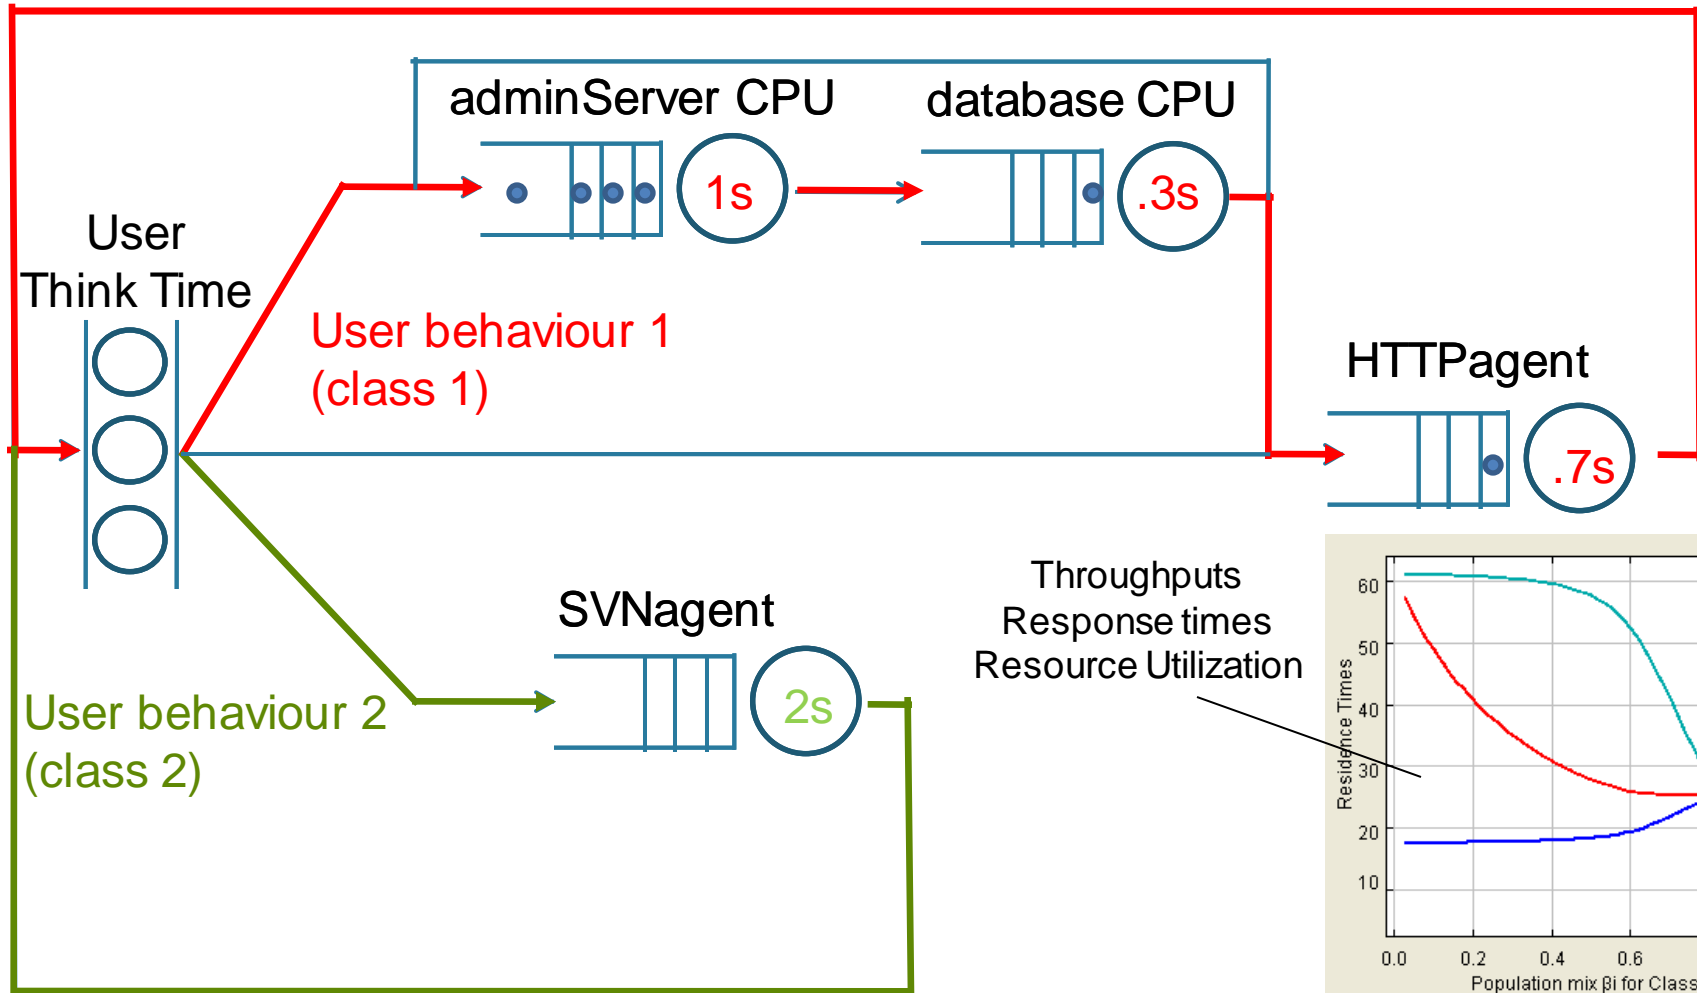
- Queueing solver
- Operational environment
- SLA percentiles

M2M Transformation CloudML Model




M2M Transformation

Queueing Network Models



Throughputs
Response times
Resource Utilization




Rest of the Talk

1. How to account for cloud-specific features in QoS modelling?
 - **Approach:** shift towards modelling the operational environment
 - **Contribution:** LINE, a solver for performance engineering in time-varying environments
2. Application: Cloud application sizing
 - Spot VM reliability model


Cloud Multi-Tenancy

- Low Contention: normal operation
- High Contention: operational slow-down caused by VM contention on the same host




Cloud Multi-Tenancy

- Low Contention: normal operation
- High Contention: operational slow-down caused by VM contention on the same host




Cloud Multi-Tenancy

- Low Contention: normal operation
- High Contention: operational slow-down caused by VM contention on the same host




Probabilistic QoS Modelling



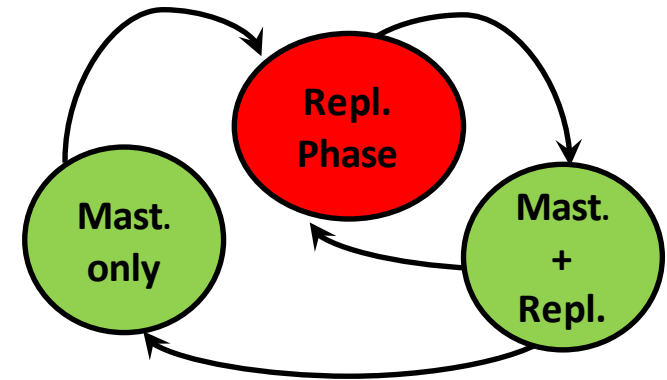
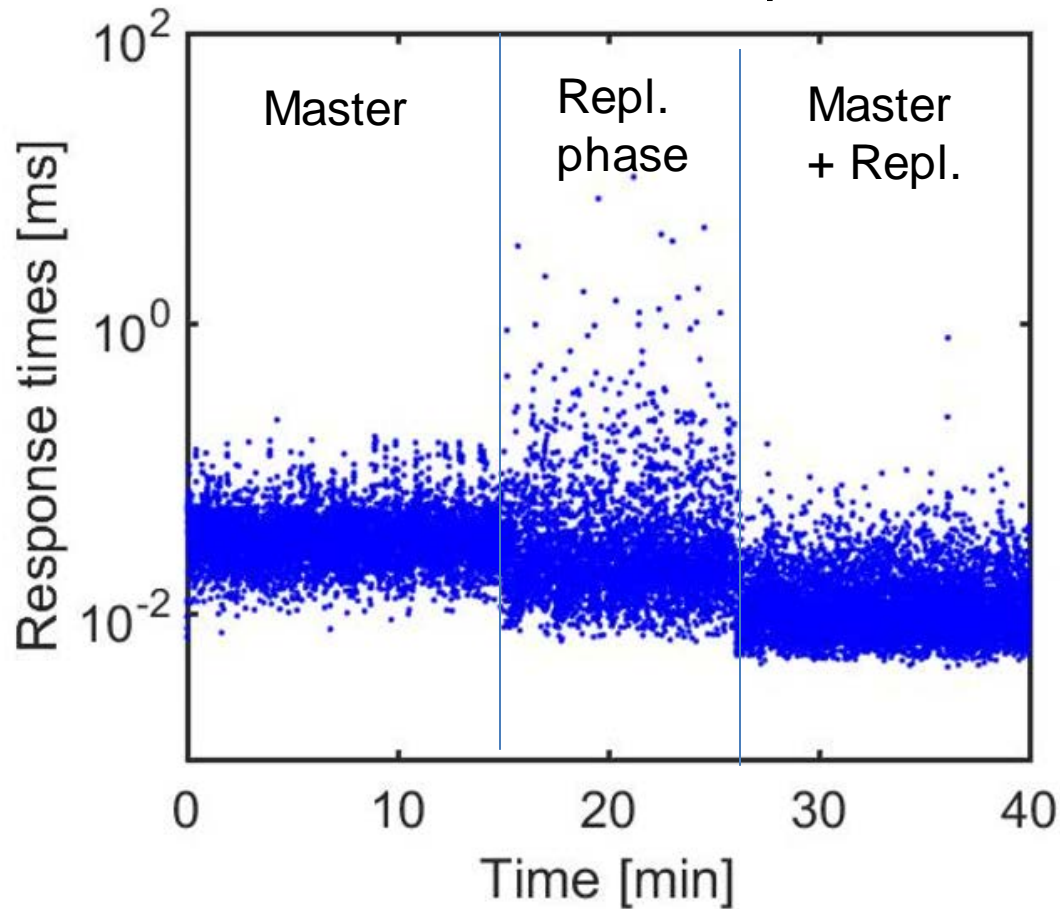
Stochastic model: CTMC

$$Q = \left[\begin{array}{c|c} Q_1^* - \alpha_{12}I & \alpha_{12}I \\ \hline \alpha_{21}I & Q_2^* - \alpha_{21}I \end{array} \right]$$




Vertical Scaling

Amazon RDS Replication



Environment Model




Spot Instances (EC2)

Spot price < Bid


Spot prices

Product: Linux/UNIX ▾ Instance type: m1.xlarge ▾ Date range: 1 day ▾ Availability zone: All zones ▾



Availability zone	Price
eu-west-1a	
eu-west-1b	
eu-west-1c	
Date	

Mouse over the graph to see prices for a specific date and time.




Environment
Model

Spot Instances (EC2)


Spot price > Bid
(VM can be reclaimed)

Spot prices

Product : Linux/UNIX ▾ Instance type: m1.xlarge ▾ Date range : 1 day ▾ Availability zone: All zones ▾


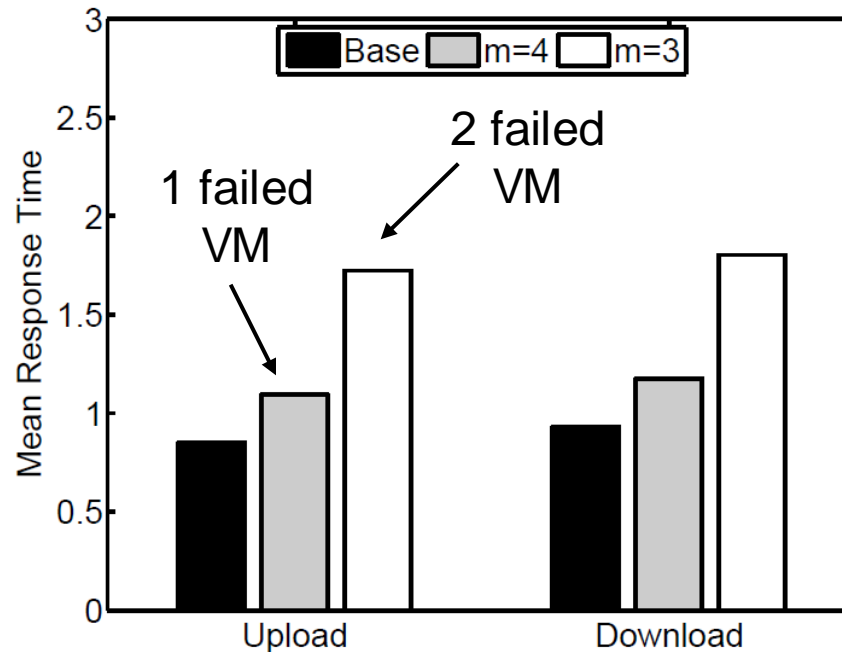


Availability zone	Price
eu-west-1a	
eu-west-1b	
eu-west-1c	
Date	




Environment Model: VM Breakdowns

Effect of VM failures **response time**:



Probabilistic QoS Modelling



Continuous-Time Markov chains (CTMCs)

$$Q = \left[\begin{array}{c|c} Q_1^* - \alpha_{12}I & \alpha_{12}I \\ \hline \alpha_{21}I & Q_2^* - \alpha_{21}I \end{array} \right]$$

Exact solution

$$\pi_1 \left(Q^{AVG} - \frac{Q_1^* Q_2^*}{\nu} \right) = 0$$

$$\pi_2 \left(Q^{AVG} - \frac{Q_2^* Q_1^*}{\nu} \right) = 0$$

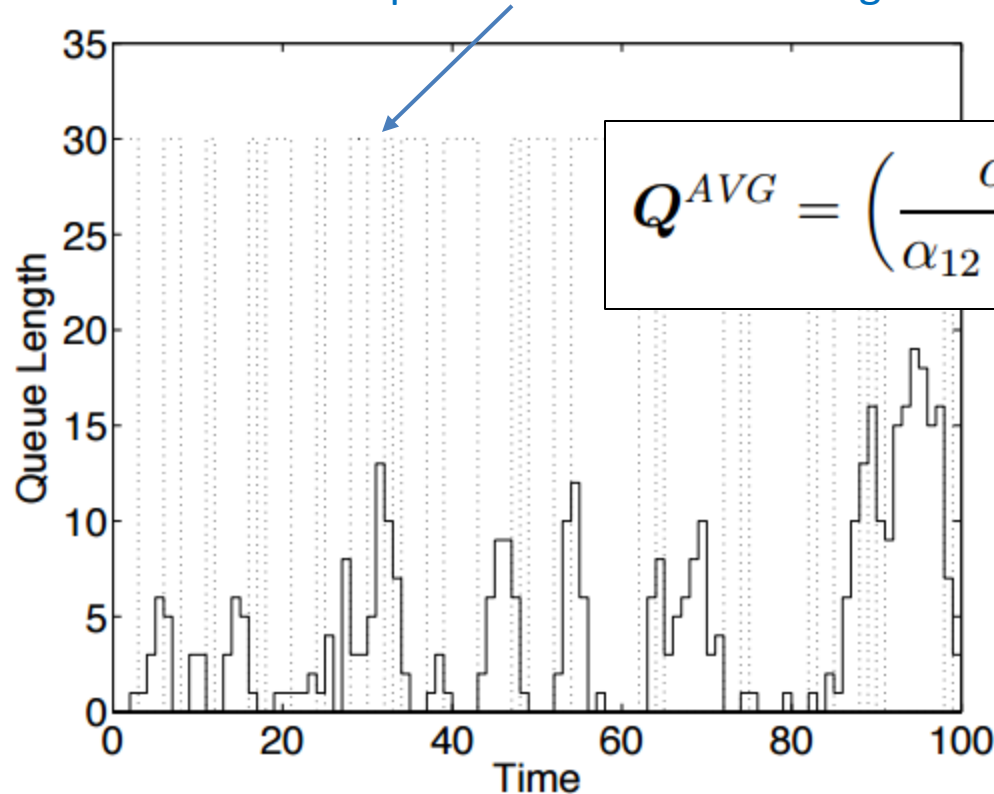
$$\nu = \alpha_{12} + \alpha_{21}$$

- Direct numerical methods seldom applicable
 - State-space explosion
 - Exponential growth with number of jobs and nodes

Average Approximation

Tandem closed queueing network (2 queues)

Limit case: frequent environment changes



Average approximation

$$Q^{AVG} = \left(\frac{\alpha_{21}}{\alpha_{12} + \alpha_{21}} \right) Q_1^* + \left(\frac{\alpha_{12}}{\alpha_{12} + \alpha_{21}} \right) Q_2^*$$

Close to exact for large jump rates


$$\nu = \alpha_{12} + \alpha_{21} \gg 0$$

(a) High ν rate (AVG approx. accurate).

Decomposition Approximation

Tandem closed queueing network (2 queues)

Limit case: rare environment changes



Decomposition method:

$$\pi_1^{DEC} Q_1^* = 0,$$

$$\pi_2^{DEC} Q_2^* = 0,$$

Results are weighted by mean
time in each environment state

Close to exact for small jump rates


$$\nu \approx 0$$

(c) Low ν rate (DEC approx. accurate).

Intermediate case

Tandem closed queueing network (2 queues)

Intermediate case



(b) Intermediate ν rate (no accurate method).

Only direct solution
methods available
(unscalable)

Contribution: LINE

- Accurate and efficient analysis of time-varying environments

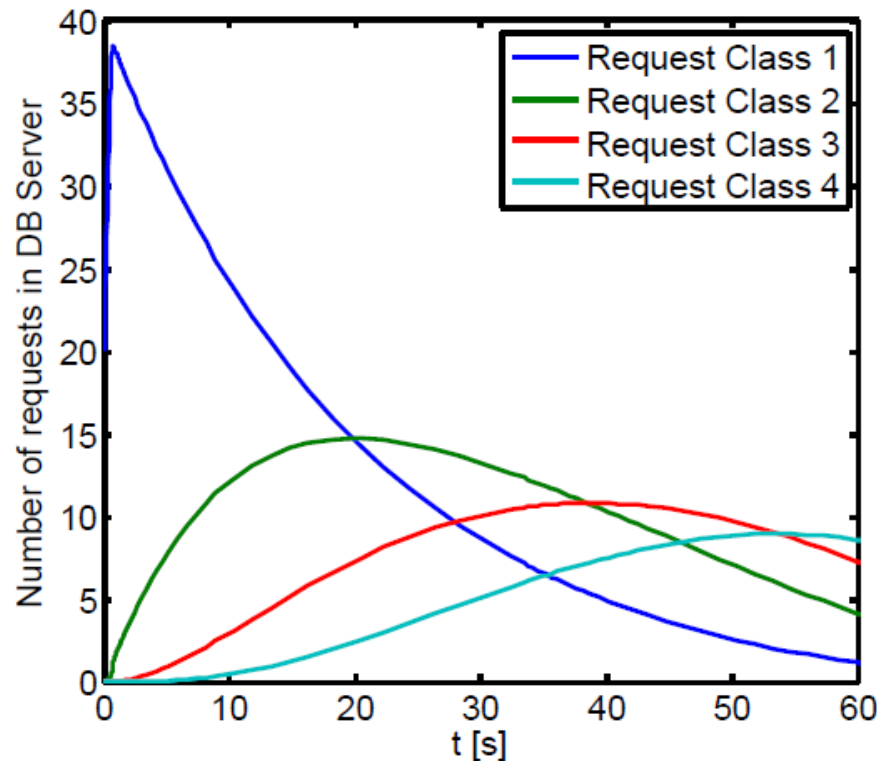
LINE Solver

- **Fluid solver** for QNs in random environments
 - Fast solution by **ordinary differential equations**
- Goes beyond mean values used in canonical solvers:
 - Transient analysis
 - Response time percentiles (e.g. SLA assessment)
 - Handling of multi-server resources (e.g., vCPUs)
 - Synchronous and asynchronous calls
 - Non-exponential service distributions (Coxian)

URL: <http://line-solver.sf.net>



Fluid Modelling

- Non-linear ODEs provide time-varying mean queue sizes
- CTMC solution converges to ODE as system size grows



Blending Method


- Leverages time-varying solutions to analyze systems operating in random environments



- **Main result:** found dynamic equilibrium between entry and exit time distributions
 - Fixed point iteration involving their Laplace transforms

Blending Method




State of node i seen at jump instants: Entry state $\bar{n}_i^e(0) = \sum_{\substack{h=1 \\ h \neq e}}^E p_{he}^{src} \alpha_h \left(\underbrace{\int_0^{+\infty} \bar{n}_i^h(t) e^{-\alpha_h t} dt}_{\text{Transient state computed by ODEs}} \right)$



Entry states at equilibrium
determined by fixed-point iteration

(b) Intermediate ν rate (no accurate method).

Example: Scaling in Amazon RDS



Avg. response time	Before Repl.	During Repl.	After Repl.
Measurement (s)	0.062	0.063	0.034
Prediction (s)	0.059	0.060	0.031
Error %	4.2%	4.4%	8.9%

Application: Sizing in the Cloud


- Provisioning of spot and on-demand resources

$m_{(i,j)}^S$ Number of spot instances in cloud i for component j
 $m_{(i,j)}^O$ Number of on-demand instances in cloud i for component j

- Provisioning based on mean:
 - Solution based on simulated annealing


$$\begin{aligned}
 \min \quad & \sum_{j=1}^2 \left(\sum_{i=1}^2 C_i^S m_{(i,j)}^S + C^O m_{(i,j)}^O \right) \\
 \text{s.t.} \quad & E[R] = \text{LINE}(m^S, m^O), \\
 & E[R] \leq R^{\max}, \\
 & m_{(i,j)}^S \in \mathbb{N}_+, \quad i = 1, 2, \quad j = 1, 2, \\
 & m_{(i,j)}^O \in \mathbb{N}_+, \quad j = 1, 2,
 \end{aligned}$$

Spot VM Reliability Model




Spot VM Reliability Model

2 spot instances in 2 clouds




Spot VM Reliability Model

Spot instances in cloud 2 are **lost**




Spot VM Reliability Model

Spot instances in cloud 2 are **lost**




Spot VM Reliability Model

“Emergency” on-demand instances **started** on cloud 1



Spot VM Reliability Model

- Response time sensitivity to time to recover a spot VM



- Provisioning based on RT percentiles:

$$R_x = \text{LINE}(m^S, m^O)$$


$$R_x \leq R_x^{\max}$$

Response time
percentile


- 1-2 orders of magnitude better than Markov ineq. & Chevyshev ineq.

Spot VM Reliability Model

- Response time sensitivity to time to recover a spot VM





Provisioning based on mean



Provisioning based on 95th-percentile

Other Applications of LINE

- Optimal bidding price for spot VMs (*Cluster Comp.*, 2016)
- Optimal cloud application refactoring (*IEEE CLOUD*, 2016)
- Inference of resource capacities in the cloud (*IEEE TSE*, 2014)
- Ongoing work: provisioning of micro-VMs
 - Environments can help model burstable performance



Questions?



g.casale@imperial.ac.uk